




Dieter Schmalstieg

The Graphics Pipeline

What do we want?

- Computer-generated imagery (CGI) of complex 3D scenes in real-time
- Computationally extremely demanding
 - Full HD at 60 Hz:
$$1920 \times 1080 \times 60 \text{ Hz} = 124 \text{ Mpx/s}$$
 - And that's just output data!
- Requires specialized hardware



FROM SD TO 8K - EXPONENTIAL PIXEL GROWTH


Image Courtesy: Rose Adler, Leighana Ginther, Jackie Osterday

Solution

Most of real-time graphics is based on

- Rasterization of graphic *primitives*

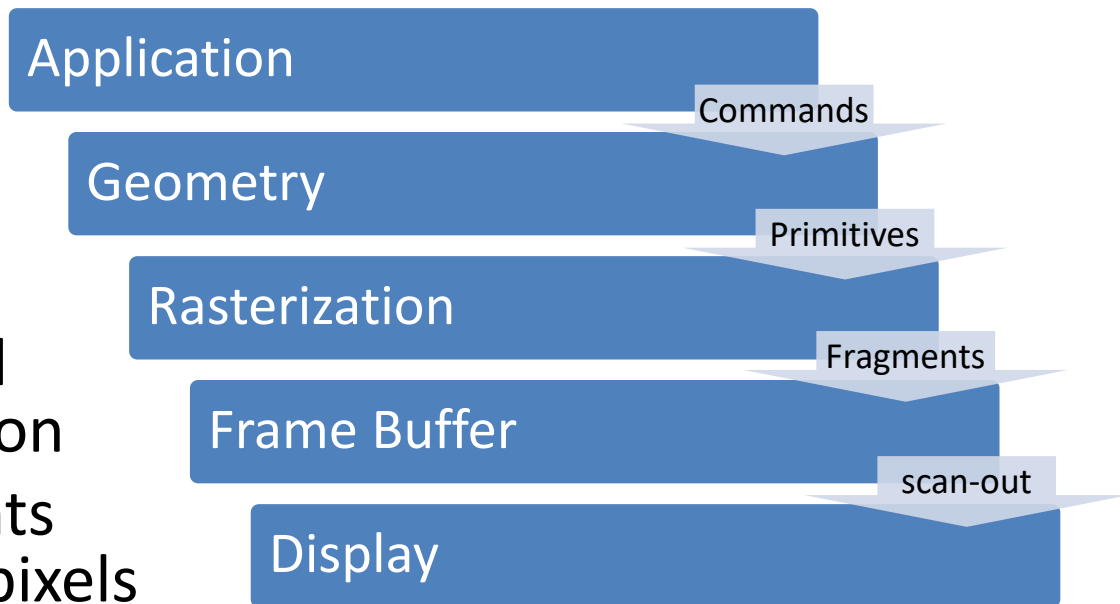
- Points
- Lines
- Triangles
- ...



- Implemented in hardware
 - *Graphics processing unit* (GPU)

The Graphics Pipeline

- High-level view:
- “Fragment”:
 - Sample produced during rasterization
 - Multiple fragments are *merged* into pixels



Application Stage

- Generate database
 - Usually only once
 - Load from disk or generate algorithmically
 - Build acceleration structures (hierarchy, ...)
- Repeat main loop
 - Input event handlers
 - Simulation → modify data structures
 - Database traversal → issue **graphics commands**
- Until exit

Graphics Commands

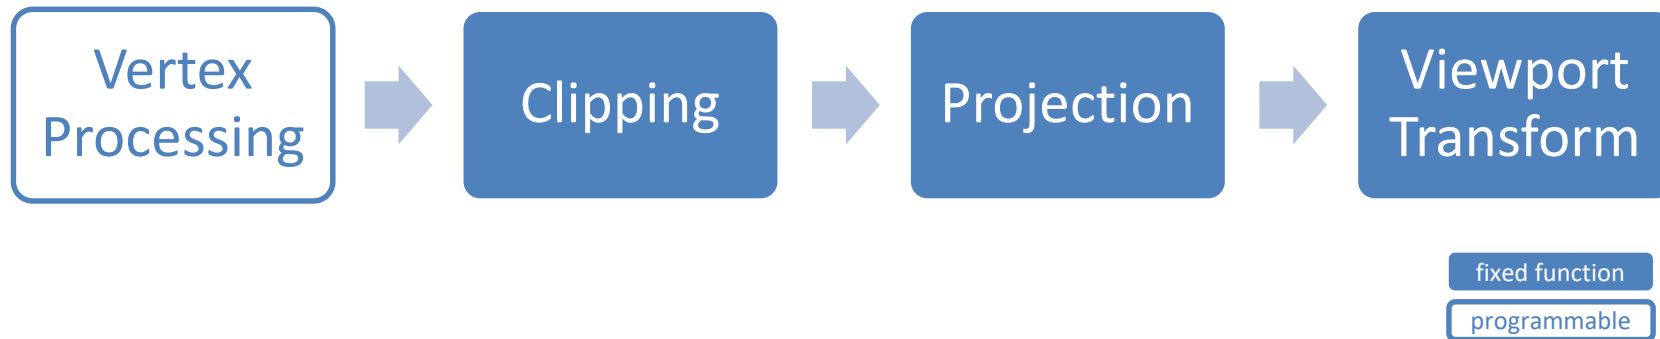
- Graphics command stream from CPU to GPU
 - Specify primitives
 - Manage resources
 - Modify GPU state

Graphics Driver

- Graphics hardware is shared resource
- Large user mode graphics driver
 - Prepares command buffers
- Graphics kernel subsystem
 - Schedule access to hardware
- Small kernel mode graphics driver
 - Submit command buffers to hardware


This is where new APIs (Apple Metal, AMD Mantle, Khronos Vulkan) try to be more efficient

Geometry Stage



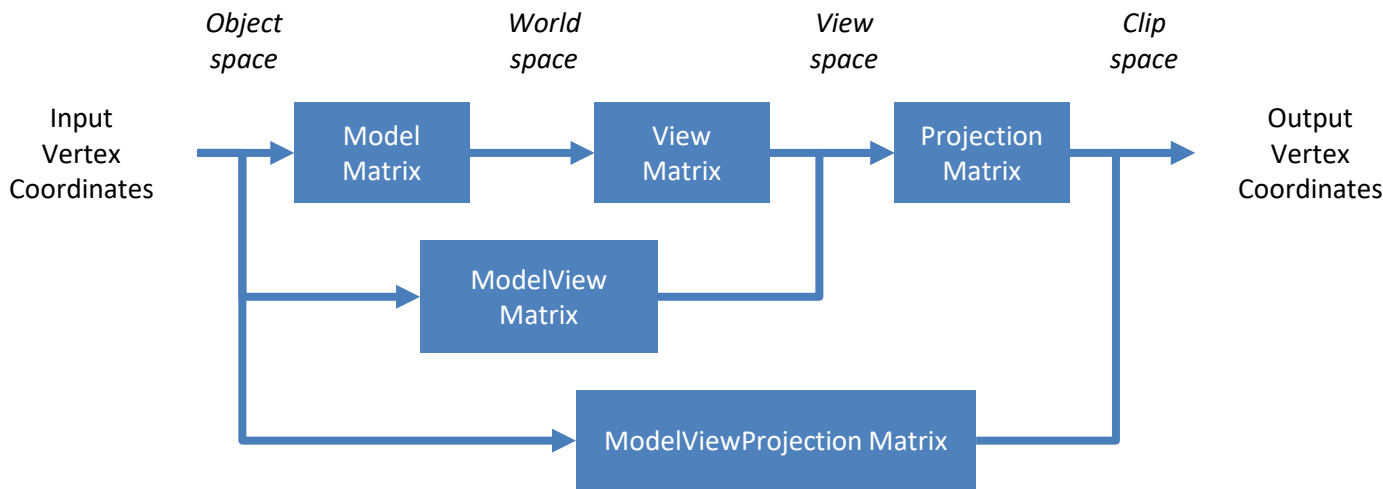
Vertex Processing

- Input vertex stream
 - Composed of arbitrary *vertex attributes* (position, color, ...)
- Is transformed into stream of vertices mapped onto the screen
 - Composed of their *clip space* coordinates and additional user-defined attributes (color, texture coordinates, ...)
 - Clip space: homogeneous coordinates
- By the *vertex shader*
 - GPU program that implements this mapping
 - Historically, “shaders” were small programs performing lighting calculations




Vertex Coordinate Transformation

Common model in rasterization-based 3D graphics

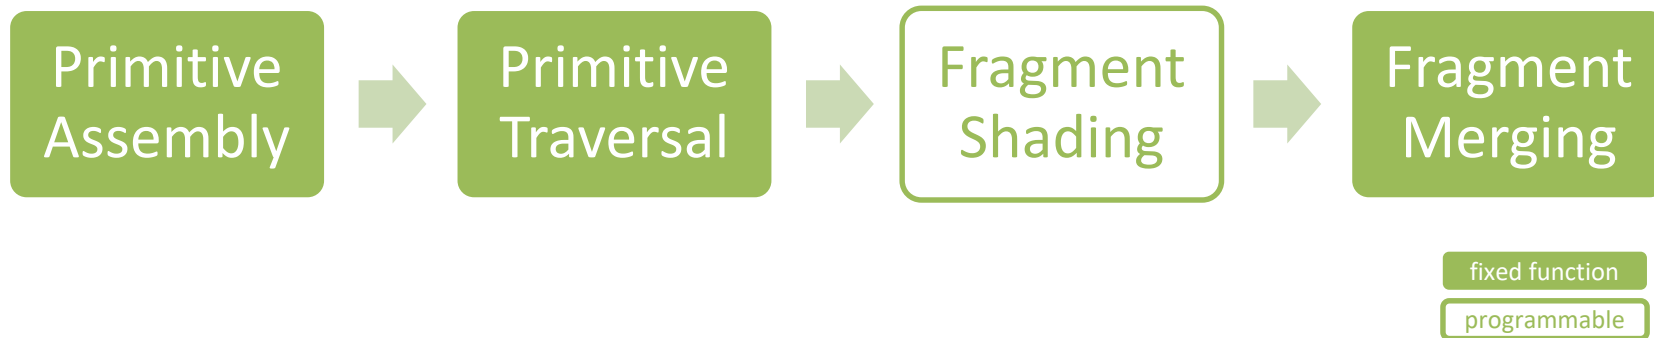


Geometry Stage Tasks

- Clipping
 - Primitives not entirely in view are clipped to avoid projection errors
- Projection
 - Projects clip space coordinates to the image plane
→ Primitives in *normalized device coordinates*
- Viewport Transform
 - Maps resolution-independent normalized device coordinates to a rectangular window in the frame buffer, the *viewport*
→ Primitives in window (pixel) coordinates




Rasterization Stage

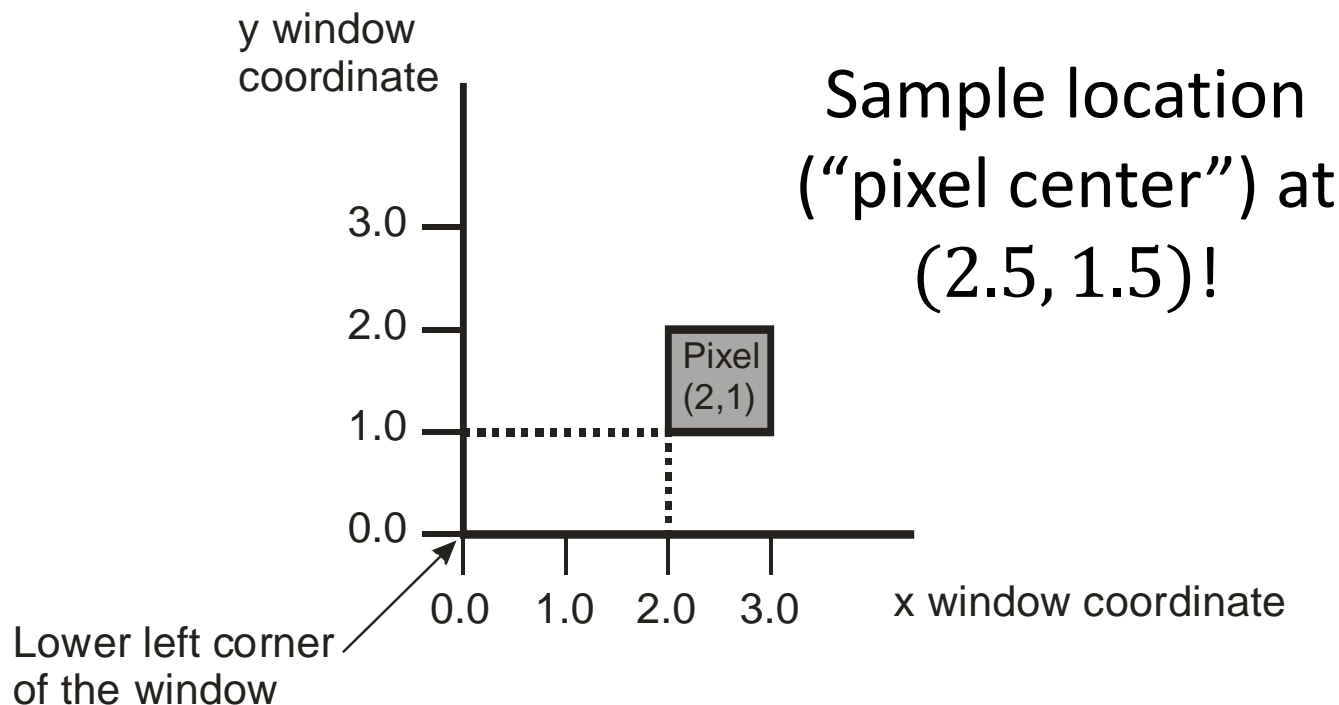


Rasterization Stage Tasks

- Primitive assembly
 - Backface culling
 - Setup primitive for traversal
- Primitive traversal (“scan conversion”)
 - Sampling (triangle → fragments)
 - Interpolation of vertex attributes (depth, color, ...)
- Fragment shading
 - Compute fragment colors
- Fragment merging
 - Compute pixel colors from fragments
 - Depth test, blending, ...




Rasterization – Coordinates




Rasterization – Rules

- Different rules apply for each primitive type
 - “Fill convention”
- Non-ambiguous!
 - Avoids artifacts
- Polygons
 - Pixel center contained in polygon
 - Pixels on edge: only one rasterized




Fragment Shading

- Given the interpolated vertex attributes
 - Output by the vertex shader
- The *fragment shader* computes color values for each fragment
 - Apply textures
 - Lighting calculations
 - ...

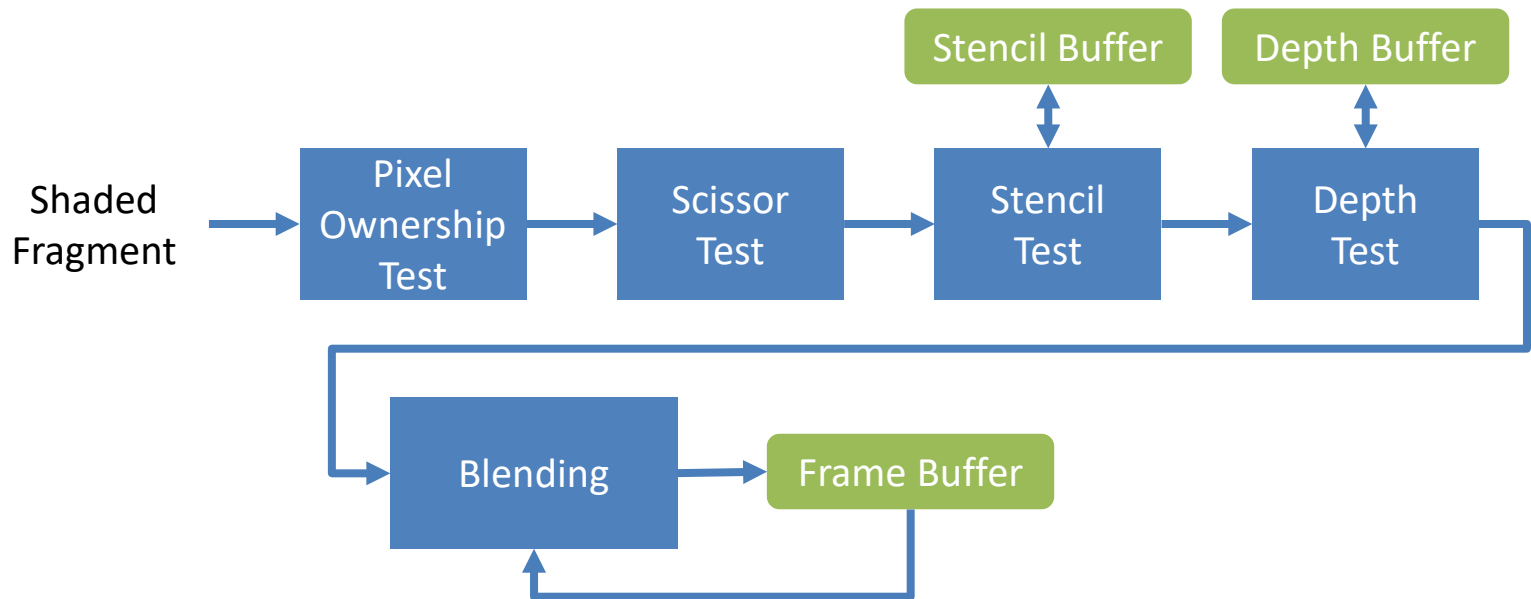


Fragment Merging

- Multiple primitives can cover the same pixel
- Their fragments need to be composed to form the final pixel values
 - Blending
 - Resolve visibility via depth buffering

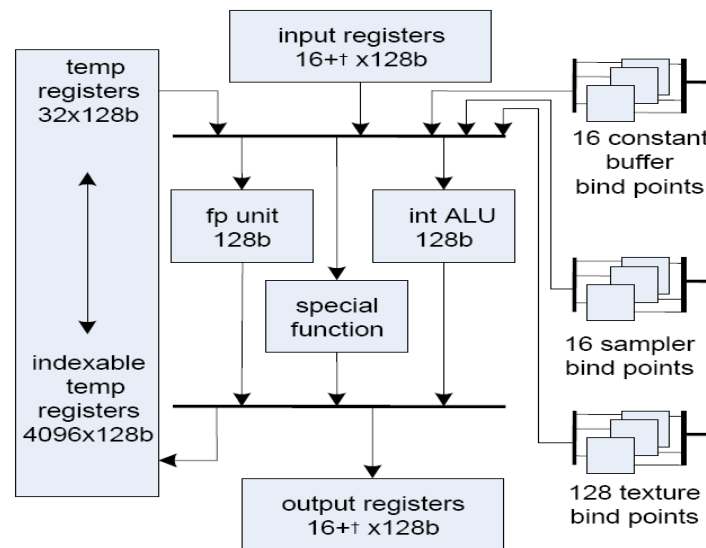


Fragment Merging Workflow



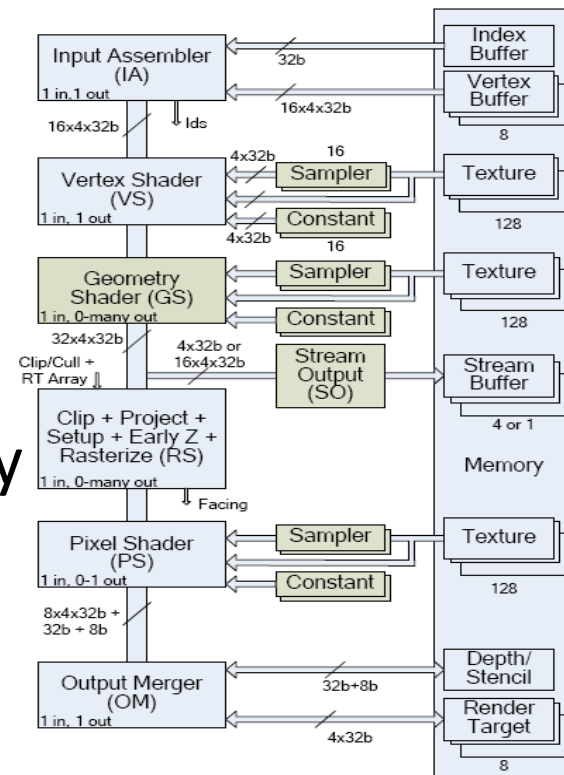
Unified Shader Model

- Same instruction set and capabilities for all shaders
- Dynamic load balancing between vertex and fragment shaders
- IEEE-754 floating point
- Enables new GPGPU languages like CUDA



Geometry Shader

- Between vertex and pixel shader
- Can generate primitives dynamically
- Procedural geometry
 - E.g., growing plants
- Geometry shader can write to memory
 - Called „stream output“
 - Enables multi-pass for geometry



Geometry Shader Examples




Figure 5: From left — render to cube map, particle system, instancing, shadow volume, displacement mapping.

- Cube Map: GS instances every triangle 6x
- Particles produced by GS as stream-out
- Plants created as parameterized instances
- Shadow volume created by GS extrusion
- Displacements created by GS

Tessellation Shader

- Since DirectX11
- New shader stage: Tessellation
- Input: low-detail mesh
- Output: high-detail mesh

Vertex Shader

Hull Shader

Tessellator

Domain Shader

Geometry Shader

Rasterizer

Pixel Shader




CryEngine3

Authoring without Tessellation

Sub-D modeling



Animation




Displacement map



Polygon mesh



Generate LODs



Authoring with Tessellation

Sub-D modeling



Animation



Displacement map




Optimally tessellated mesh



Display Stage

- Gamma correction
- Historically: digital to analog conversion
- Today: digital scan-out, HDMI encryption?

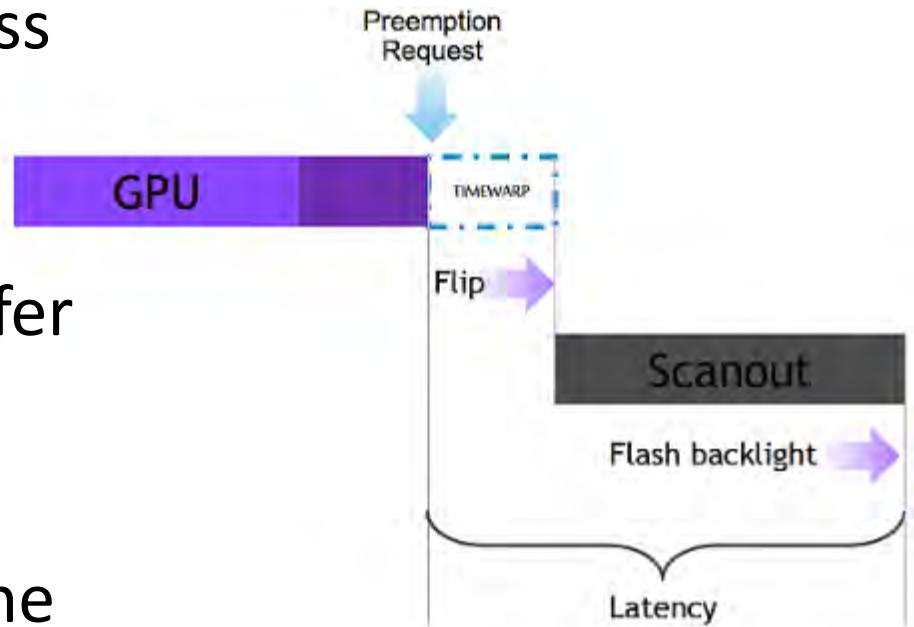


Display Format

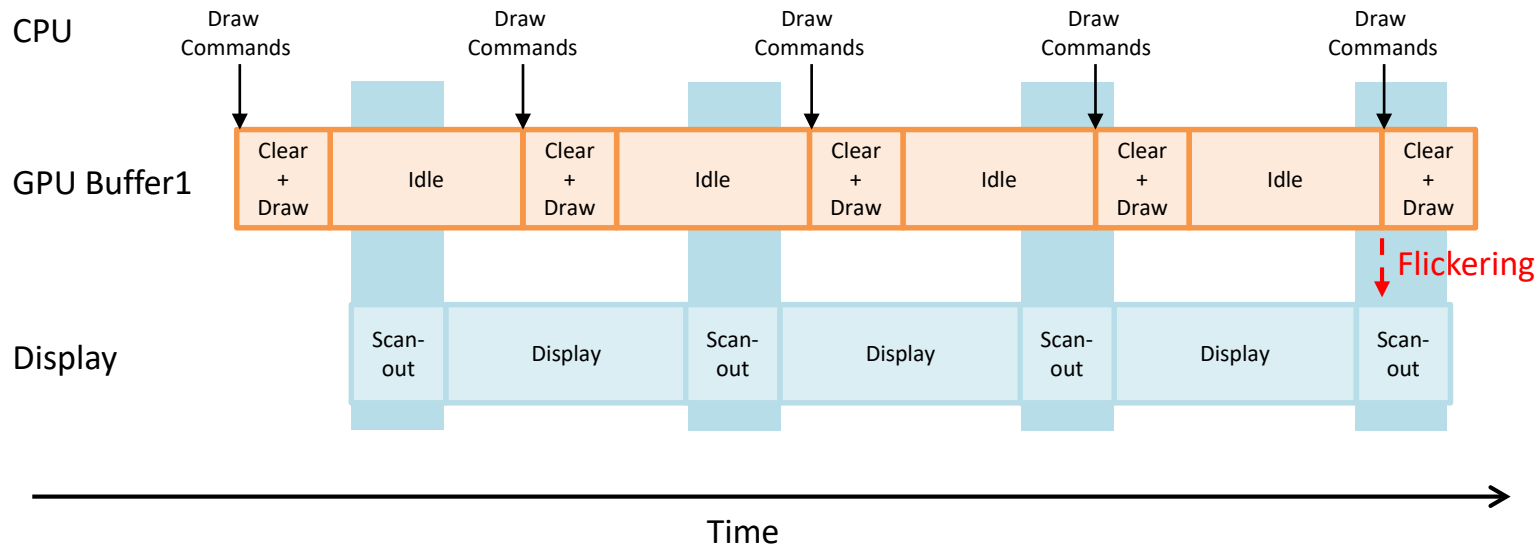
- Frame buffer pixel format:
RGBA vs. index (obsolete)
- Bits: 16, 32, 64, 128 bit floating point, ...
- Double buffering for smooth animation
- Quad-buffering for stereo graphics
- Overlays (extra bitplanes)
- Auxilliary buffers: alpha, stencil, depth

What is Display Synchronization?

- Need to synchronize access to frame buffer between GPU and display
- Cannot change frame buffer during display scan-out
- Need proper buffering in the whole graphics pipeline

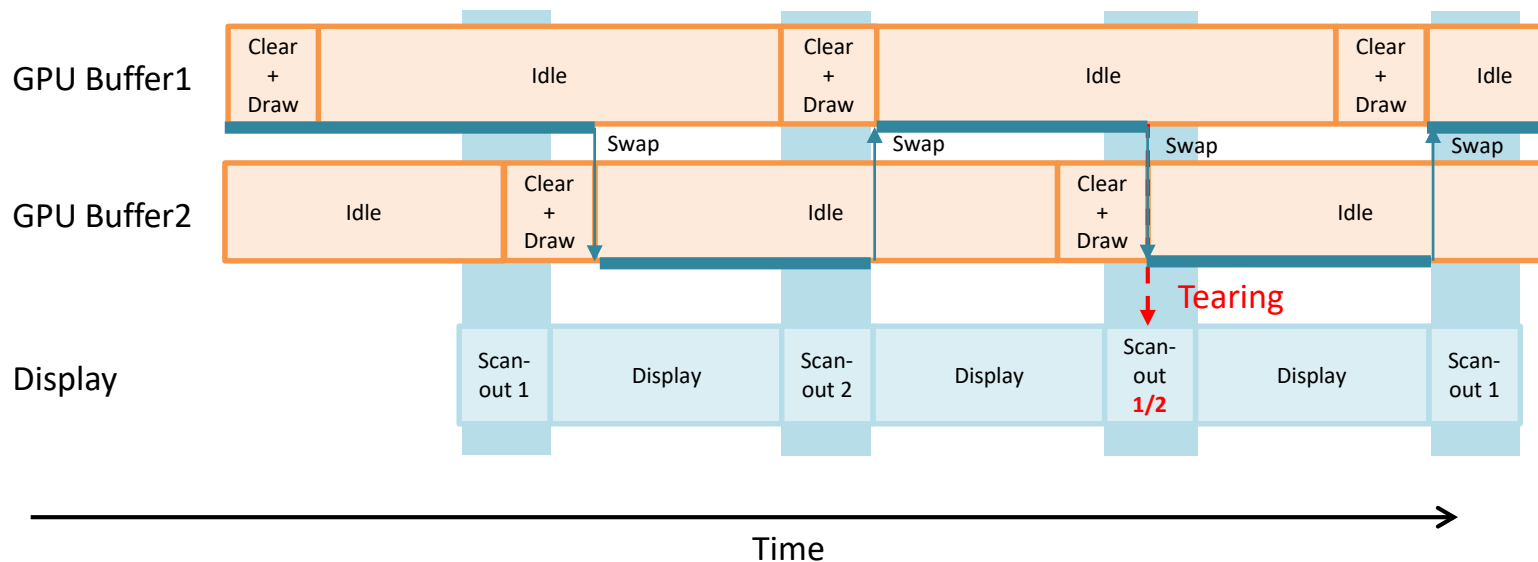



Single Buffering



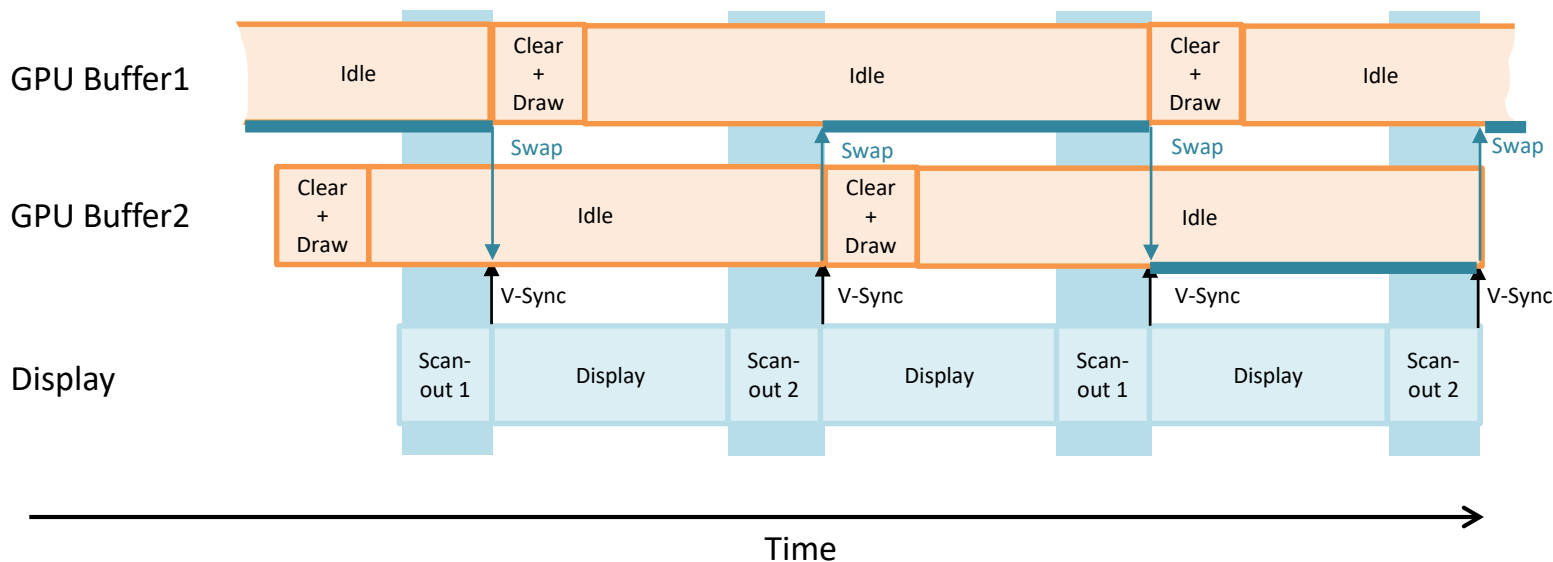
Double Buffering without V-Sync

- Front buffer used for scan-out GPU → display
- SwapBuffers() changes front and back buffer
- No flickering, but **tearing** artefacts



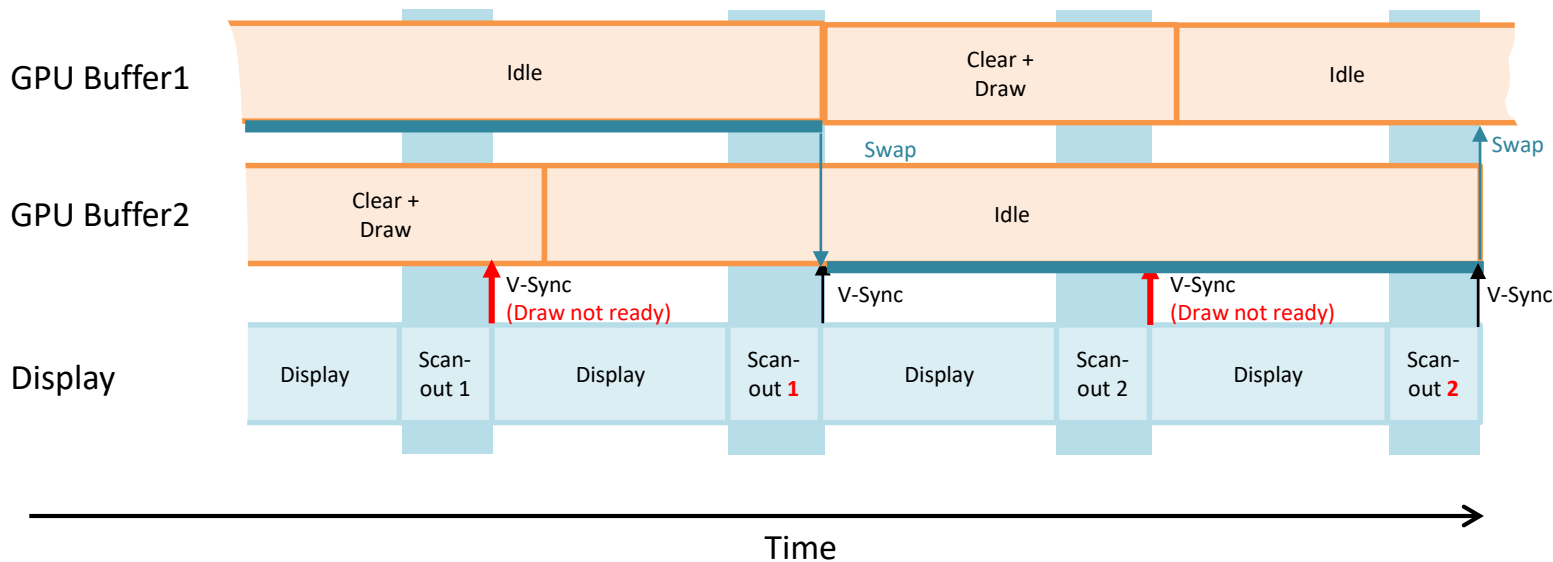
Double Buffering with V-Sync, Fast

- V-Sync means the display is done with frame scan-out from GPU
- When rendering is **fast**, the frame rate is limited by display rate
- **Additional latency** of max. one frame time



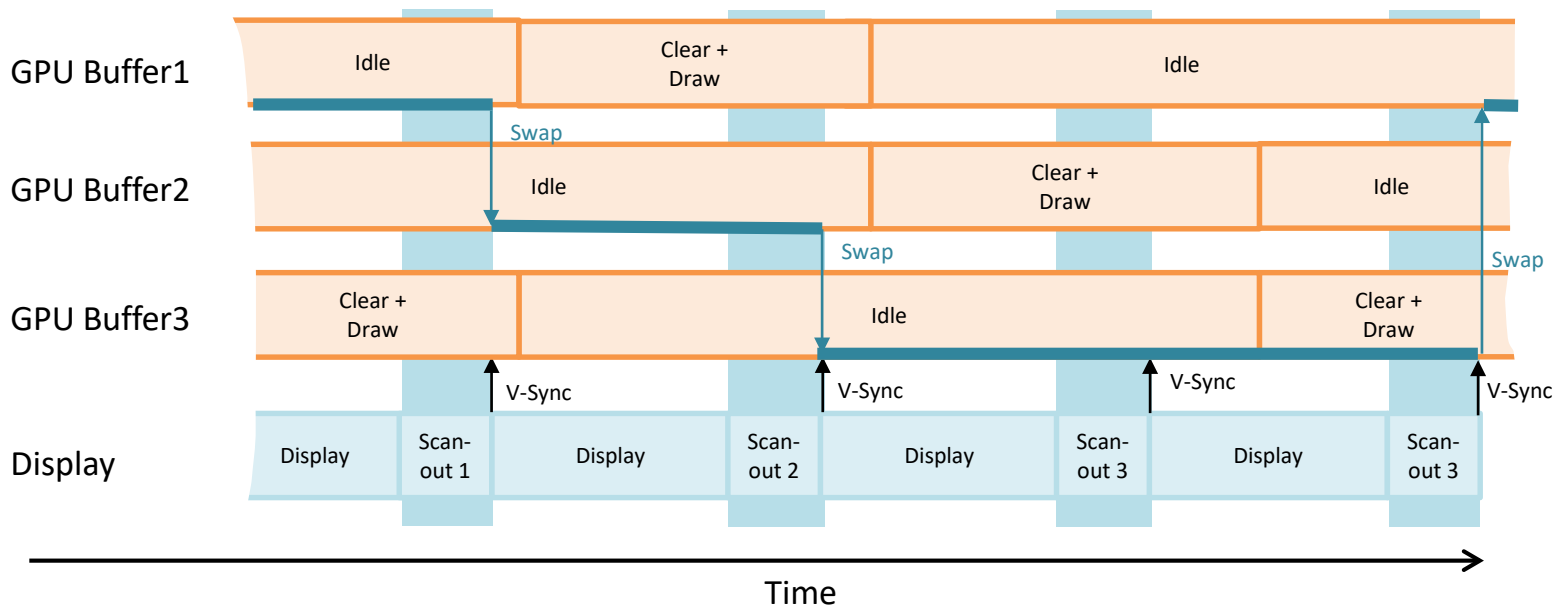
Double Buffering with V-Sync, Slow

- When rendering is **slow**, frame rates can **only be integer fractions of display rate**
- Display rate 60Hz → frame rates 60, 30, 20, 15, ... (but not 59) possible
- Adaptive V-Sync (NVIDIA) turns off V-Sync automatically if rendering is slow



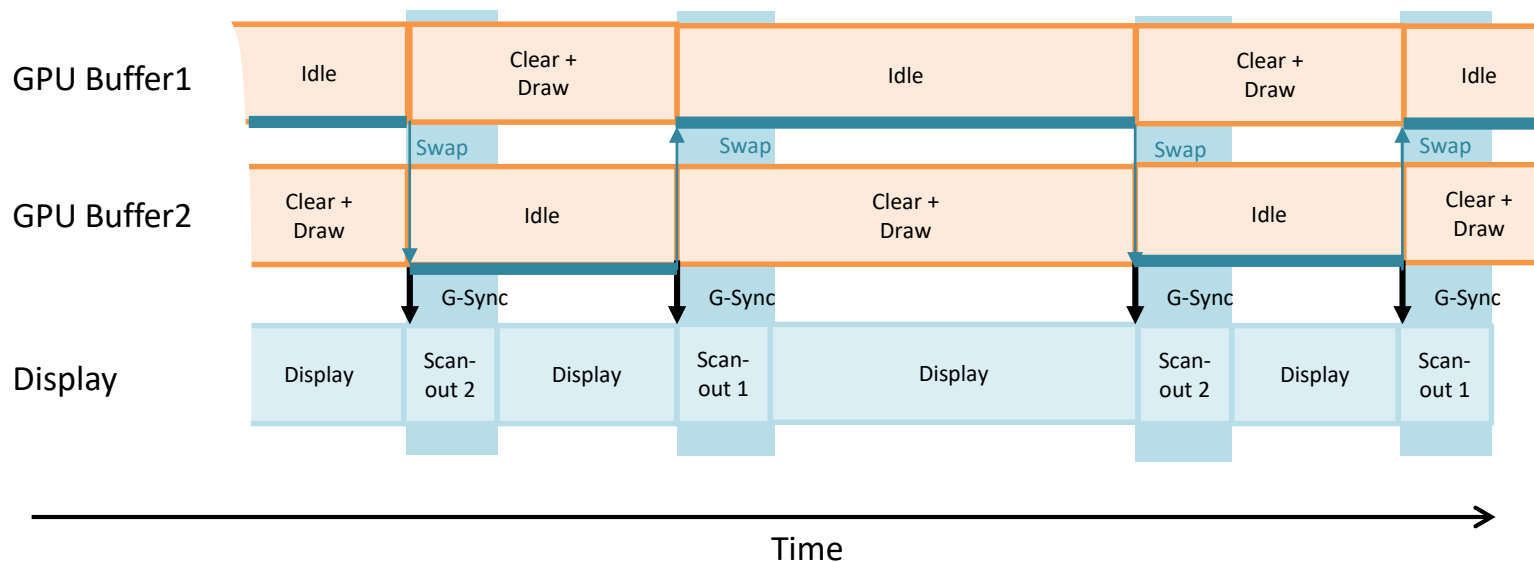
Triple Buffering with V-Sync

- Triple buffering only makes sense together with V-Sync
- When rendering is **slow**, third buffer allows **continuous drawing**
- When rendering is **fast**, two back-buffers can be alternated (**wasted** frames!)



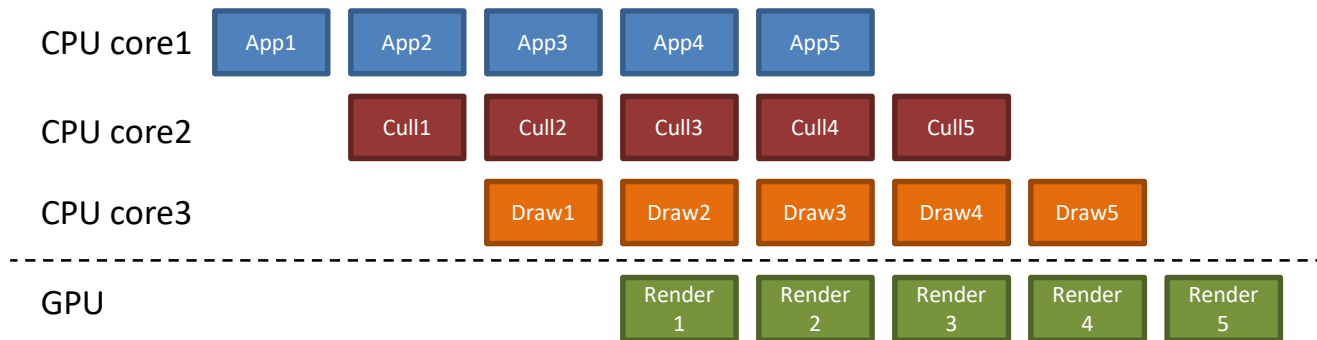
G-Sync / FreeSync

- Display scan-out can be triggered by GPU
- Requires additional display hardware feature
- G-Sync (NVIDIA), FreeSync (AMD)



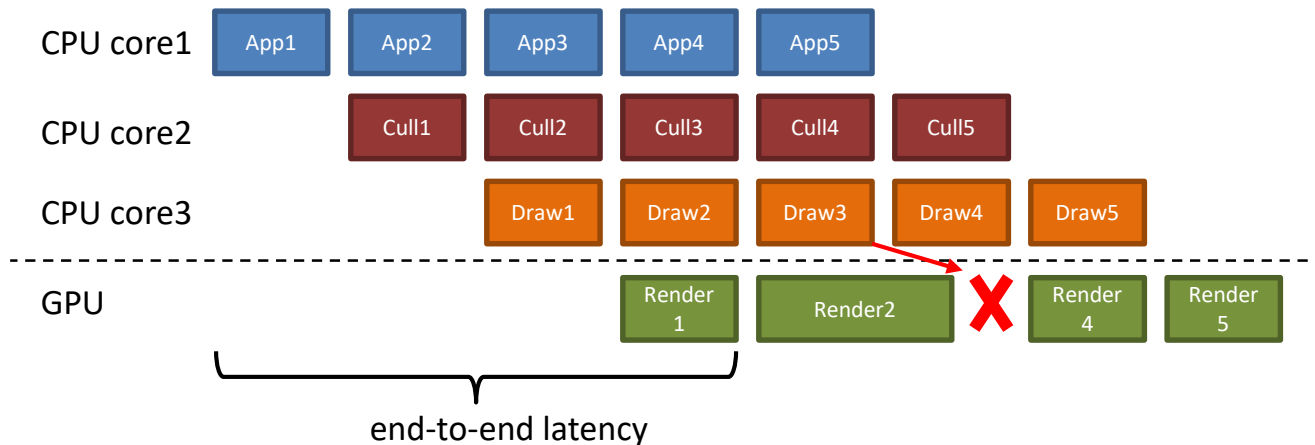
Multi-Threaded Rendering Pipeline

- App stage: simulate 3D world
- Cull stage: determine object in view frustum
- Draw stage: issue OpenGL commands to driver (includes optimizations such as mode sorting)
- Everything must work at target frame rate!




Minimizing Pipeline Latency

- Always aim for minimal latency
- Fixed size buffer from stage to stage
- Never wait for (downstream) consumer!



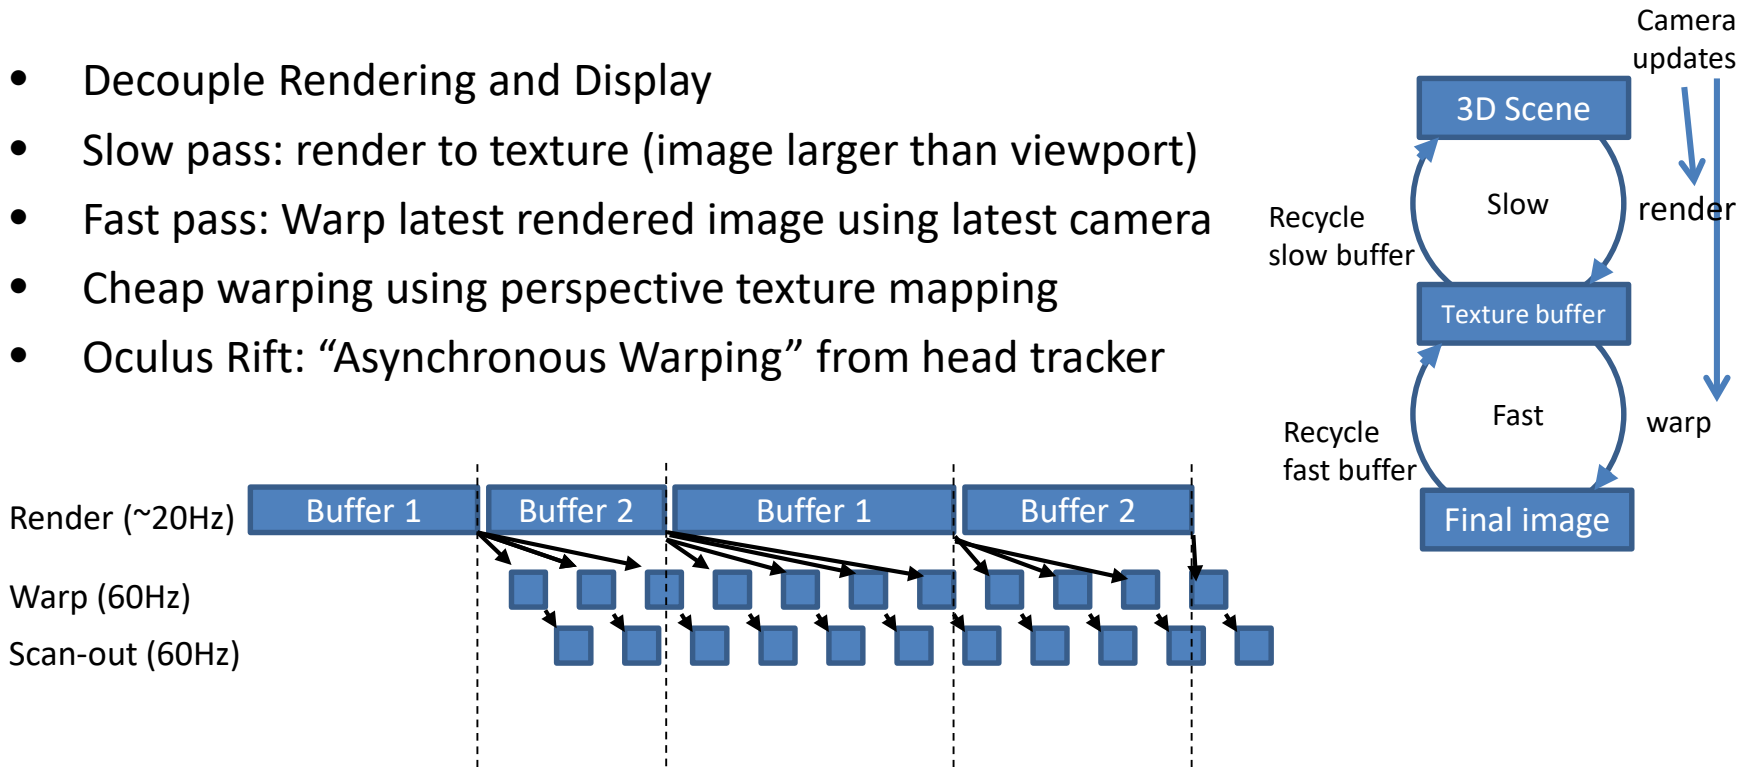
Warping

- Additional warping stage
- Pass 1: Render to Texture
- Pass 2: Perspective texture mapping with new camera parameters



Multi-Framerate Rendering

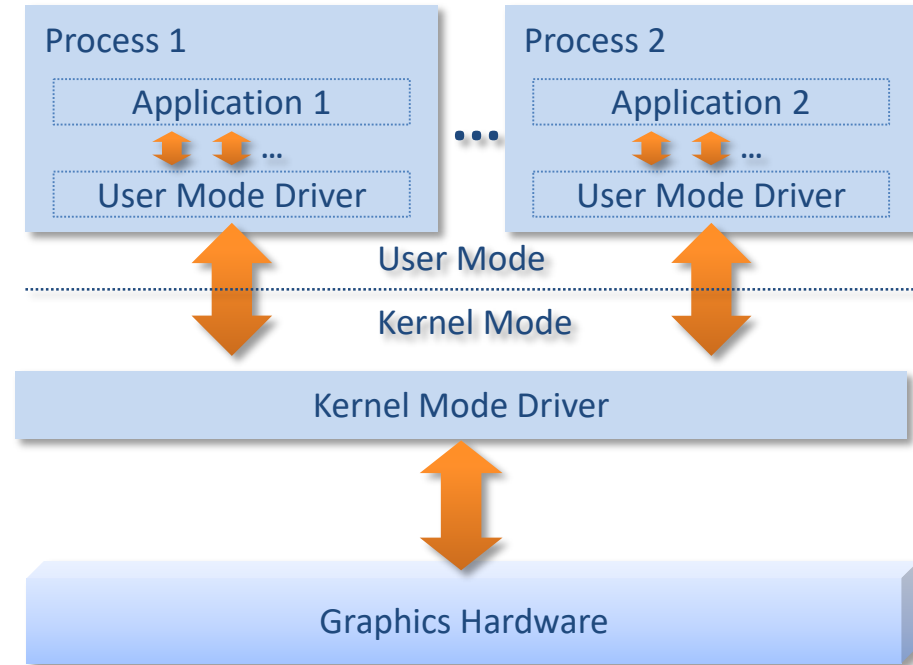
- Decouple Rendering and Display
- Slow pass: render to texture (image larger than viewport)
- Fast pass: Warp latest rendered image using latest camera
- Cheap warping using perspective texture mapping
- Oculus Rift: “Asynchronous Warping” from head tracker



NEXT-GEN GRAPHICS API

Graphics Driver Architecture

- User mode graphics driver
 - Minimize number of mode switches
 - Translation of graphics commands to instructions for the hardware
 - Batching, optimization, validation
 - Fine grained memory management
- Kernel mode graphics driver
 - Schedule access to hardware
 - Microkernel pattern, stability
 - Coarse grained memory management
 - Submits command buffers to GPU



What is a GPU Pipeline?

- Every application sees a virtualized GPU
 - Via **rendering context** object provided by driver to app
- Inside a context, one must use commands to configure GPU **pipeline state** before actual rendering
- Pipeline state consists of
 - State variables: blend, depth, culling, etc.
 - Shaders
 - Layout: how to map settings/bindings at each stage's shader

Traditional Pipeline Design

- Traditional = OpenGL (any) or DirectX (< version 12)
- Problem 1
 - Incremental configuration of state is costly
 - Submitting configuration changes to driver requires immediate validation, conversion, buffering
- Problem 2
 - Pipeline state *not* made explicit in rendering context
 - Drivers must have per-game optimizations built in

Next-Gen Pipeline Design

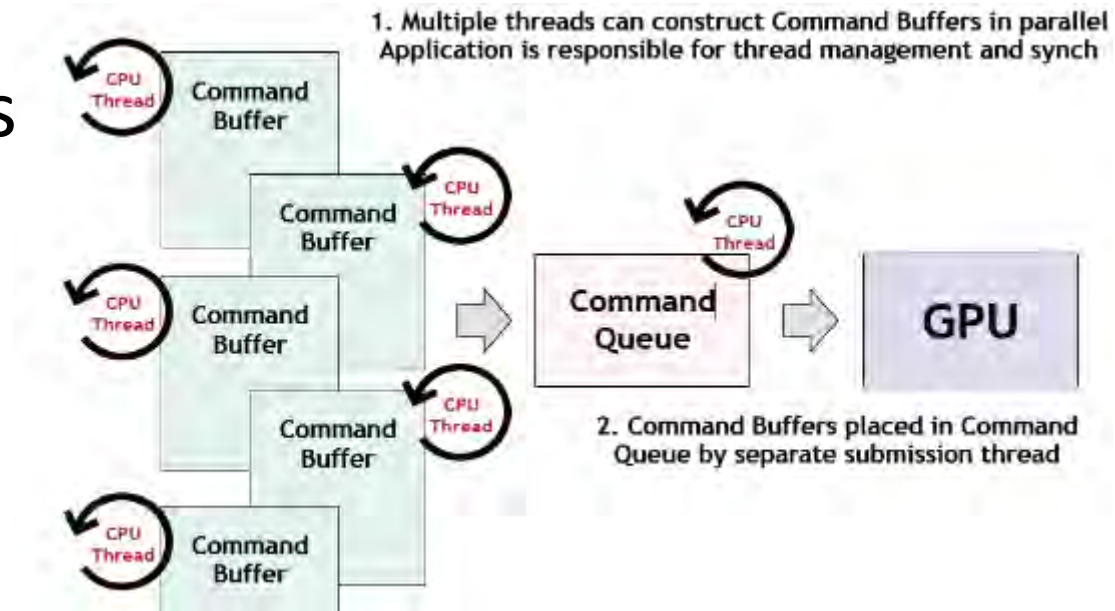
- New API: DirectX 12, Vulkan (OpenGL successor)
- Better fit for modern hardware (including mobile)
- Make (CPU driver side of) pipeline programmable
 - Multi-threading at your own risk
- Split rendering context into command buffers and dispatch queues
- Make pipeline state and render passes explicit
- Low overhead

Command Buffers

- Commands collected in command buffers
 - Optimize and validate command buffers during building, not during submission
 - Yields *immutable*, re-useable pipeline state configurations
 - Selected pipeline state variables can be declared *mutable*
- Multi-pass rendering just *switches* pipeline states
- Build *many* command buffers from *many* threads
- Can use *synchronization* primitives across buffers
 - Event, barrier, semaphore, fence

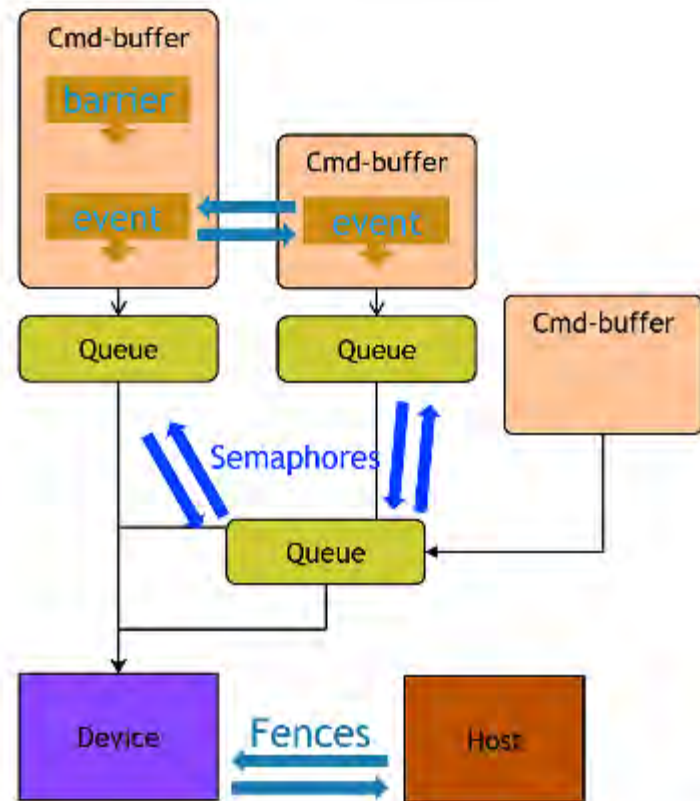
Queues

- Queues replace traditional contexts
- Insert command buffer into queue to schedule it



Synchronization

- **Events/barriers** synchronize within a queue's buffer(s)
- **Semaphores** synchronize across queues
- **Fences** synchronize between GPU and CPU




Vulkan Shaders

- Compilation separated into front-end and back-end
- Front-end: GLSL, OpenCL, or new shader languages
- Back-end = SPIR-V
 - Standard Portable Intermediate Representation
 - Byte-code derived from LLVM
- Games ship with SPIR-V, not shader source
- Vendor neutral: JIT compile to NVIDIA, AMD...
 - HLSL uses intermediate code too, but not vendor-neutral




ARM Vulkan Benchmark

- ARM Cortex A-15/7, Mali T-628 MP6
- 1000 meshes, 3 materials
- 79% less CPU



Microsoft Asteroids Demo



Context Switch Problem

- Multiple queues submit commands
- Every queue has different context (pipeline state)
- Only one queue may submit
- Without preemption
 - New jobs must wait
- With preemption
 - New high-priority job can interrupt running job
 - But context switch introduces overhead



Async Compute

- GPU with async compute
 - Multiple queues can submit commands
 - No overhead
- Particularly important when switching between graphics and compute shaders



AsyncCompute.mp4